

Large Scale Distributed Systems


José Orlando Pereira

Departamento de Informática
Universidade do Minho




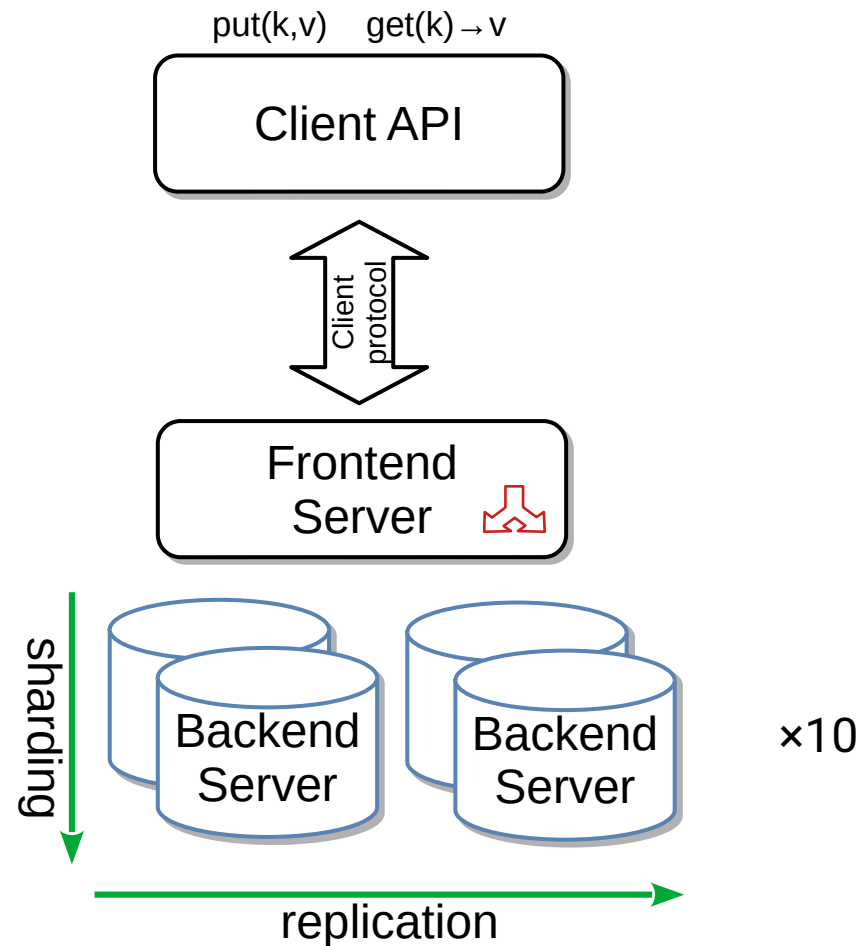
Key-Value Store

- Simple data model:
 - $\text{Map}\langle K, V \rangle$
- Simple interface:
 - $\text{get}(k) \rightarrow v$
 - $\text{put}(k, v)$
- Avoids session state in server:
 - No multi-item transactions
 - No long lived operations



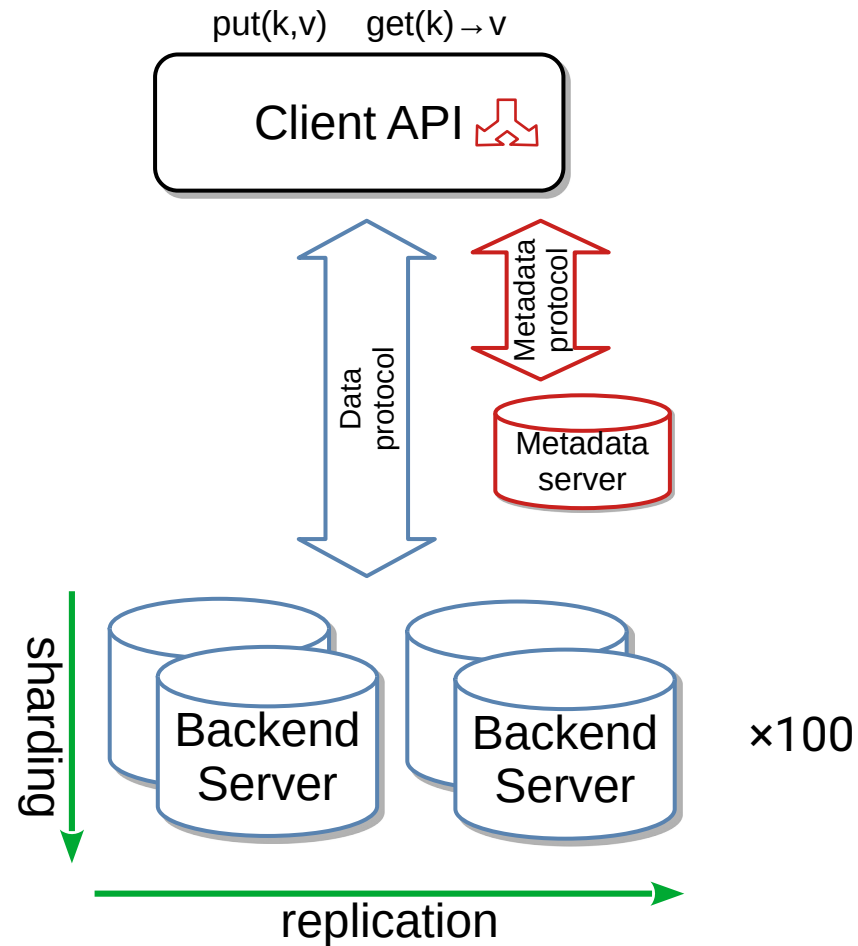
Centralized

- Replication for availability
- Sharding for scale-out
- Centralized architecture:
 - Frontend server with metadata store
 - Backend servers with data stores
- Frontend is a bottleneck (latency and bandwidth)
- Example:  **mongoDB**






Decentralized data

- Separate data and metadata servers and protocols
 - Client-side caching of metadata
- Avoids data bottleneck
- Examples:



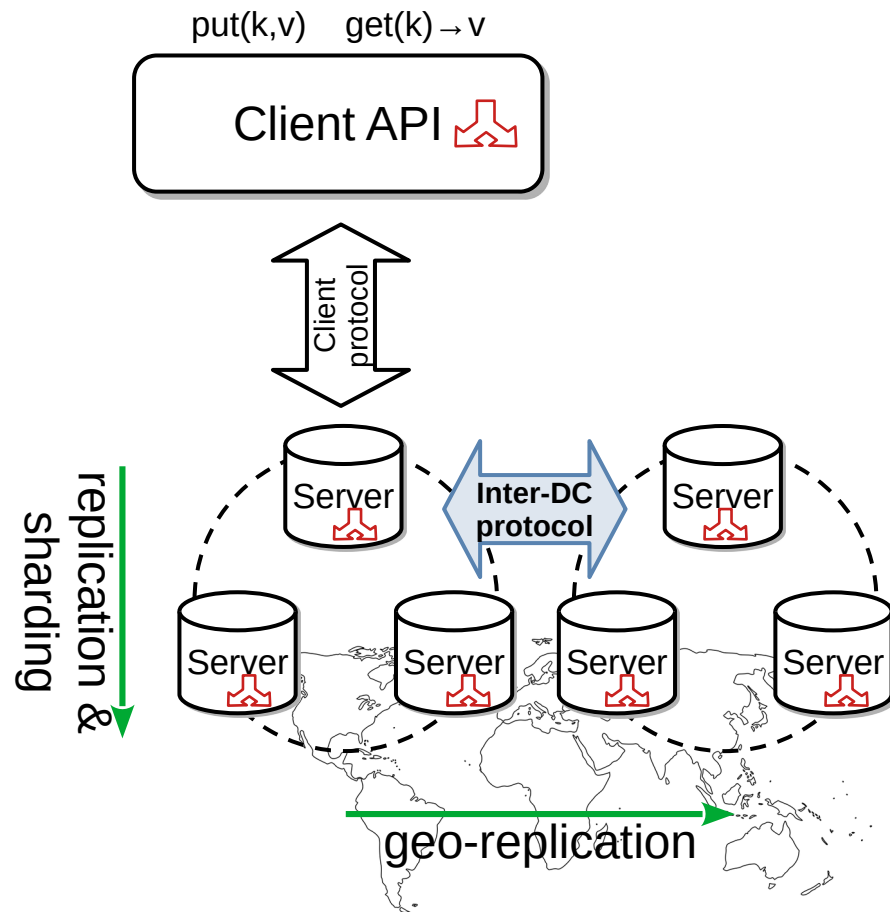
Fully decentralized

- Consistent hashing
- Epidemic dissemination
- Examples:
 - Dynamo
(not DynamoDB!) 
 -  Apache CASSANDRA

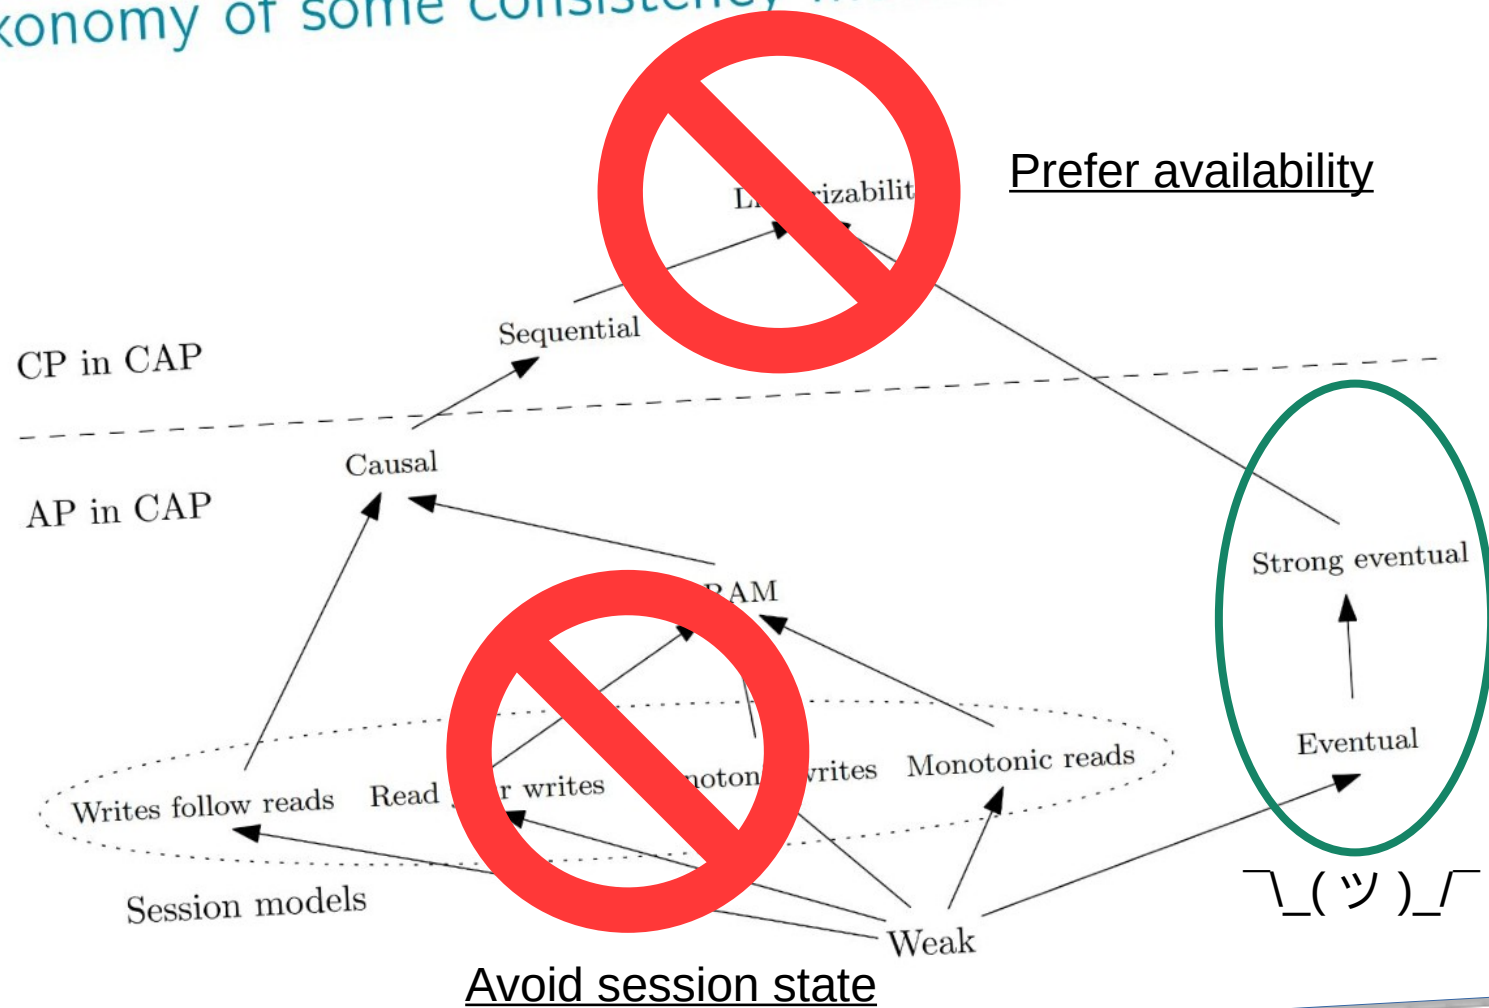


Geo-replication

- Multi-data center replication
- Main challenge: How to shield clients from Inter-DC protocol?
 - Latency
 - Availability




A taxonomy of some consistency models




Local writes

- Writes are issued to one or a few servers in the local data center
- Writes are asynchronously propagated to other data centers



Convergence

- Simple reconciliation rule: Last Writer Wins (LWW)
 - Attach a timestamp to each data item
 - On conflict, keep the item with the latest timestamp



Convergence


- Generating the timestamp:
 - Physical: when the client is a Web application and the state partitioned by user
 - Logical?
- LWW Register is actually a State-based CRDT... and KVS servers could support a variety of CRDTs
 - Limits the range of possible data structures and operations
 - Pushes complexity / policy / computation into the server
 - Not worth it if divergence is rare (client affinity + network stability)

Convergence

- Idea: Delegate reconciliation to the client application
 - Simple LWW (Cassandra)
 - Other State-based CRDTs (Cure)
 - Custom code (Bayou)
- Servers cannot callback into the client application
- How to achieve this?


Convergence

- Keep list of conflicting values in each key and return them to the application: $\text{get}(k) \rightarrow \langle v_1, \dots, v_n \rangle$
- Let the application merge them and replace the list with a single value



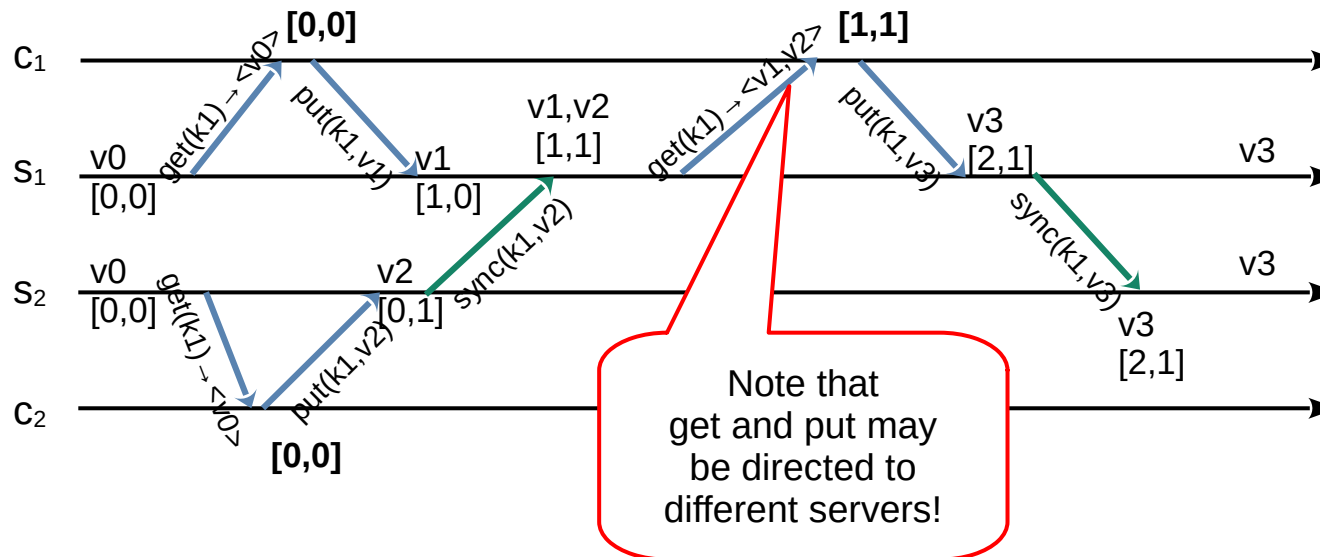
Ordering propagation requests

- Keep a vector timestamp with each data item
 - An element for each server
- When merging concurrent timestamps, keep distinct values

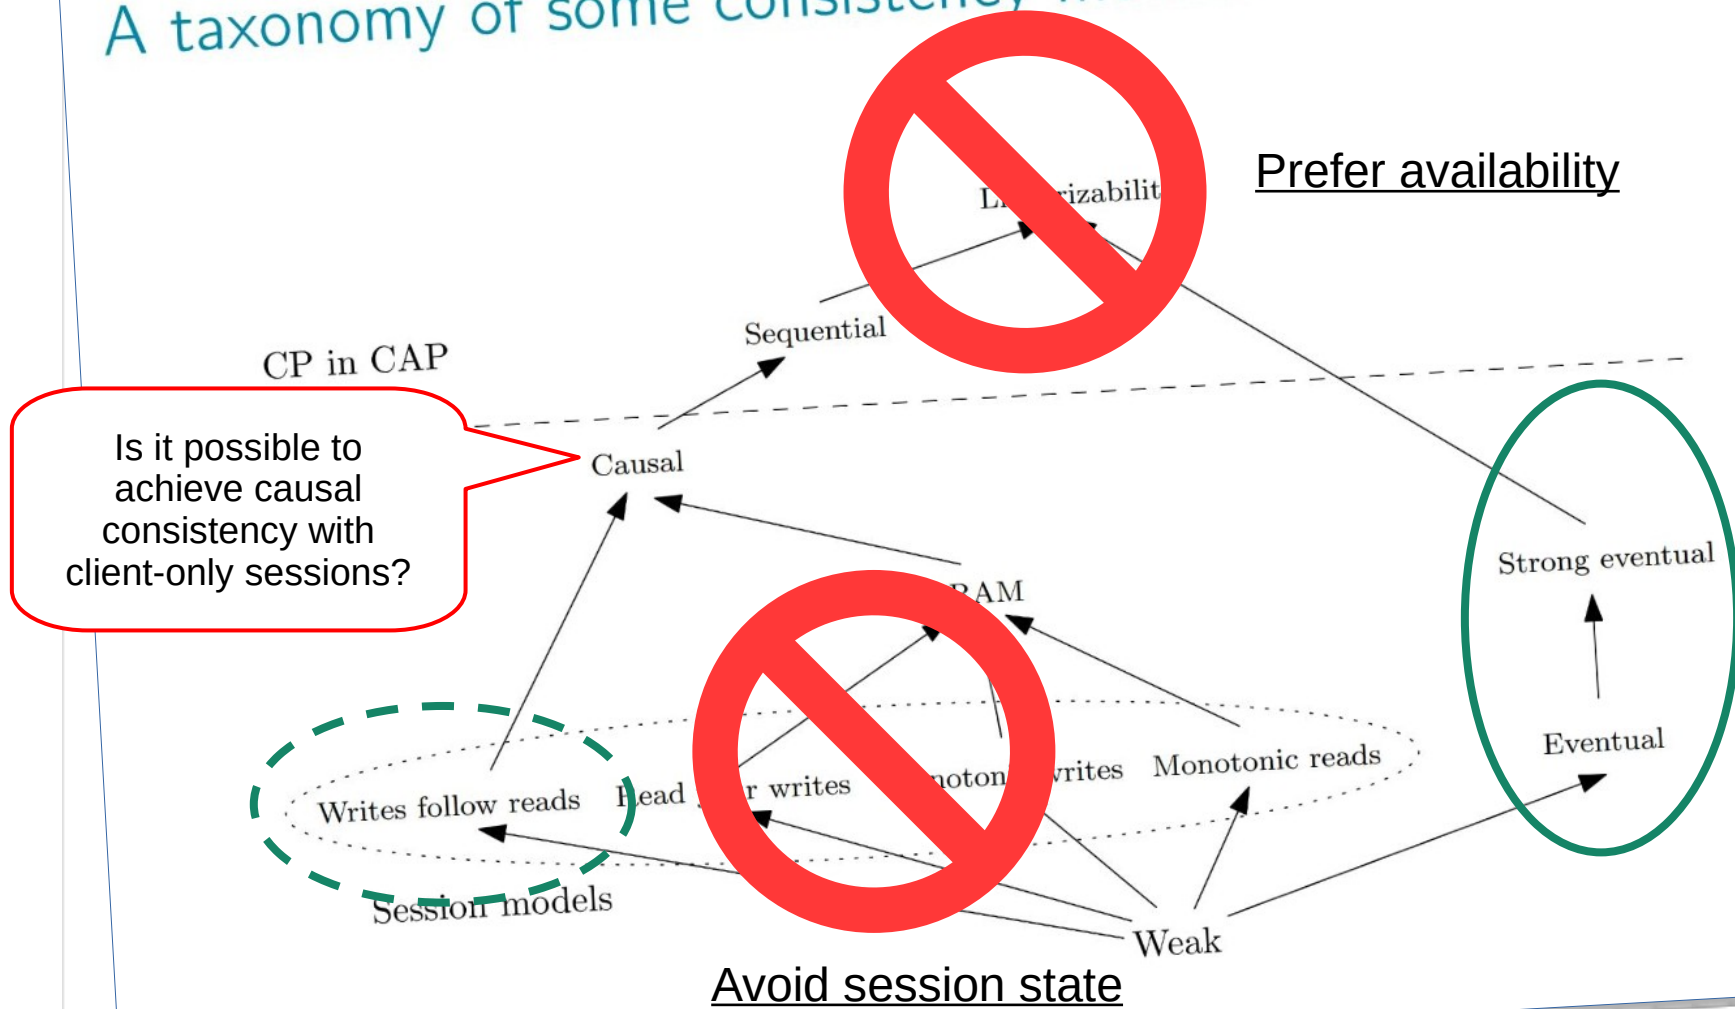


Ordering write requests

- Client requests are anonymous: How to tell what the client has read before?
 - Force the client to read before writing
 - Make it keep track of context




A taxonomy of some consistency models



Causal KVS (COPS)

- Client API:
 - `createCtx()`→ctx
 - `get(k,ctx)`→v
 - `put(k,v,ctx)`
 - `deleteCtx(ctx)`
- Servers keep a scalar timestamp for each (k,v) item (including server id in lower bits)
- Server protocol:
 - `get(k)`→(v,vers)
 - `put_after(k,v,deps,vers=0)`→vers
 - `dep_check(k,vers)`



get operation is trivial

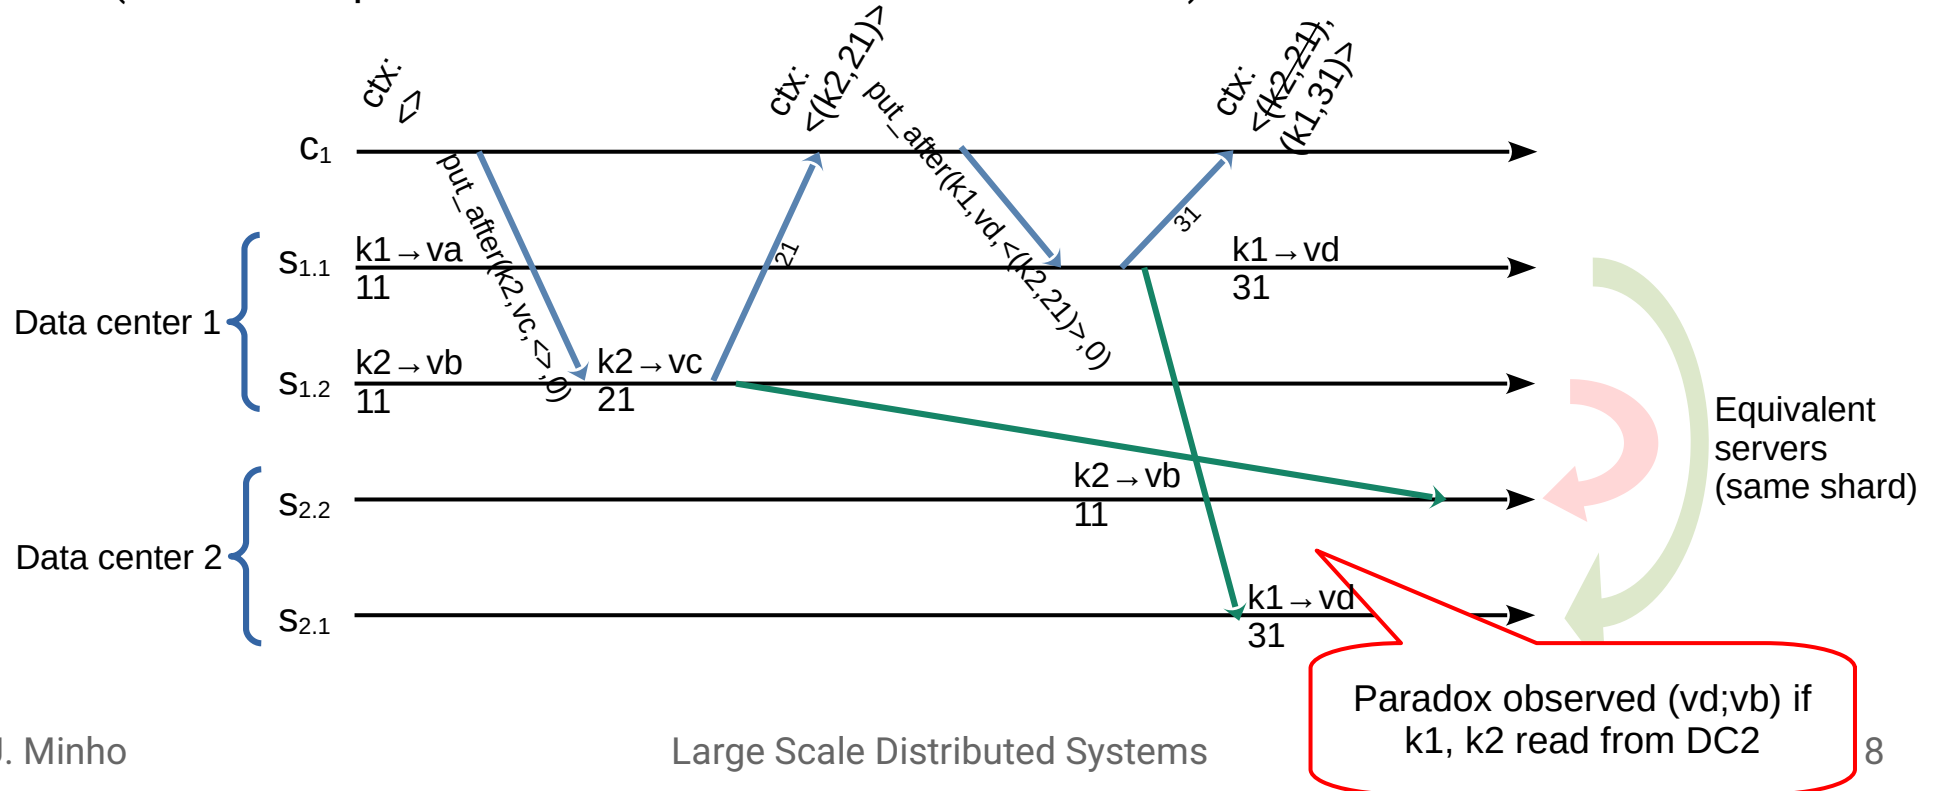
Managing client context

- Client context needs to keep nearest dependencies
- Simple (approximate) strategy:
 - After a $\text{put}(k, \dots) \rightarrow \text{vers}$ operation:
 - Clear all dependencies
 - Insert (k, vers)
 - After a $\text{get}(k, \dots) \rightarrow (v, \text{vers})$ operation:
 - If (k, \dots) in context, update to (k, vers)
 - If not, insert (k, vers)

(It is approximate because it does not recognize dependencies between separately read keys and keeps them both)


Write operation

- Synchronous write to server in local data center:
 - append dependencies in local context to request
 - wait for write to commit before returning local clock
 (it is now impossible to read an older k in the same DC)



Write propagation

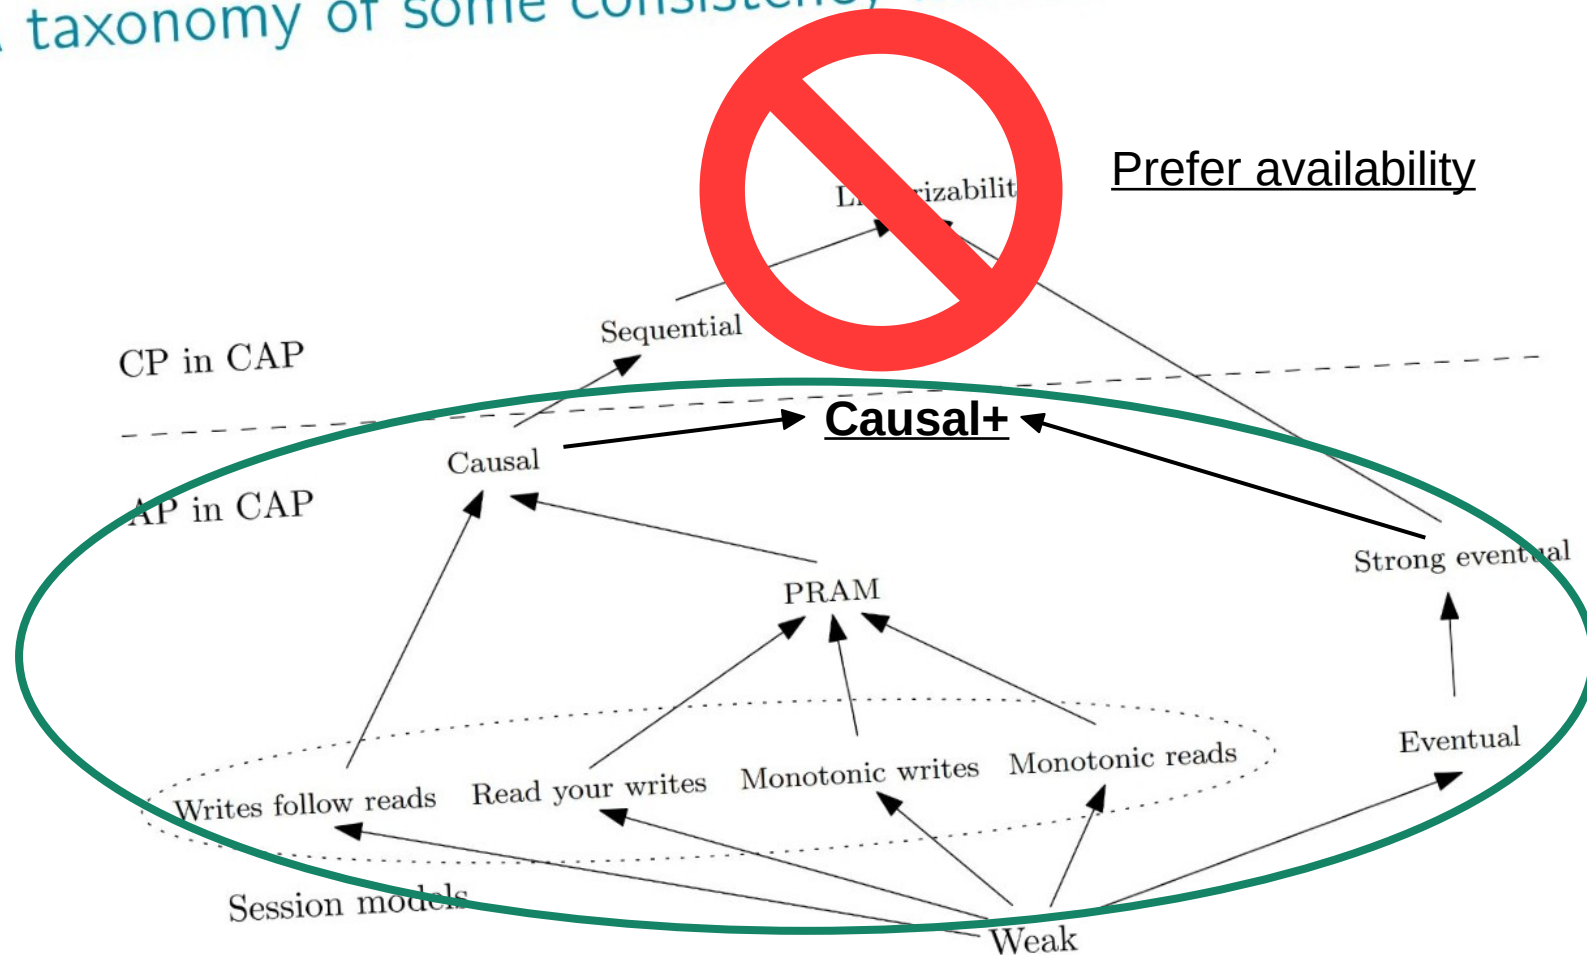
- Propagation is not ordered, as it is potentially from different data centers
- A value is committed remotely only after waiting for dependencies with the dep_check operation of corresponding servers



Consistency and convergence

- WFR: reads kept in local context and used in write
 - RYW:
 - MR:
 - MW:
- } synchronous write and client affinity
- Strong Eventual:
 - Logical timestamps including global server mean that there are no ties between concurrent updates to same item
 - Last Writer Wins when propagating to remote DC

A taxonomy of some consistency models



References

- G. DeCandia et al., “**Dynamo: Amazon’s highly available key-value store**,” Oper. Syst. Rev., vol. 41, no. 6, pp. 205–220, Oct. 2007
<https://doi.org/10.1145/1323293.1294281>
- W. Lloyd, M. J. Freedman, M. Kaminsky, and D. G. Andersen, “**Don’t settle for eventual: Scalable causal consistency for wide-area storage with COPS**,” in Proc. 23rd ACM Symposium on Operating Systems Principles, 2011
<https://www.cs.cmu.edu/~dga/papers/cops-sosp2011.pdf>